

Chapter 7: Memory Organization and Assembly Language Programming

Objectives

When you are finished with this lesson you will be able to

- describe how a typical memory system is organized in terms of memory address modes;
- describe the relationship between a computer's instruction set architecture and its assembly language instruction set; and
- use simple addressing modes to write a simple assembly language program.

Introduction

This lesson will begin our transition from hardware designers back to software engineers. We'll take what we've learned so far about the behavior of the hardware and see how it relates to the instruction set architecture (ISA) of a typical microprocessor. Our study of the architecture will be from the perspective of a software developer who needs to understand the architecture in order to use it to their best advantage. In that sense, our study of assembly language will be a metaphor for the study of the architecture of the computer. This perspective is quite different from that of someone who wants to be able to design computer hardware. Our focus throughout this book has been on the understanding of the hardware and architectural issues in order to make our software development complement the design of the underlying hardware.

We're first going to investigate the architecture of the Motorola 68000 processor. The 68K ISA is a very mature architecture, having first come out in the early 1980's. It is a fair question to ask, "Why am I learning such an old computer architecture? Isn't it obsolete?" The answer is a resounding "No!" The 68K architecture is one of the most popular computer architectures of all time and derivatives of the 68K are still being designed into new products all the time. Palm PDA's still use a 68K derivative processor. Also, a completely new processor from Motorola, the ColdFIRE family, uses the 68K architecture, and today it is one of the most popular processors in use. For example, the ColdFIRE is used in many inkjet printers.

In subsequent chapters we will also look at two other popular architectures, the Intel X86 processor family and the ARM family. With these three microprocessor architectures under our belts we'll then be familiar with the three most popular computer architectures in the world today. Much of what we'll be discussing as we begin our investigation will be common to all three processor families, so studying one architecture is as good as another when it comes to gaining an understanding of basic principles.

Another reason for starting with the 68000 is that the architecture lends itself to the learning process. The memory addressing modes are straight-forward, and from a

software developer's point of view, the linkages to high level languages are easy to understand.

We'll start by looking at the memory model again and use that as a jumping-off point for our study of how the computer actually reads and writes to memory. From a hardware perspective we already know that because we just got finished studying it, but we'll now look at it from the ISA perspective. Along the way, we'll see why this sometimes strange way of looking at memory is actually very important from the point of view of higher level languages, such as C and C++. Ladies and gentlemen, start your engines. . . .

Memory Storage Conventions

When you start looking at a computer from the architectural level you soon realize that the relationship between the processor and memory is one of the key factors in defining the behavior and operational characteristics of the machine. So much of how a computer actually accomplishes its programming tasks revolves around its interface to memory. As you'll soon see when we begin to program in assembly language, most of the work that the computer does involves moving data in and out of memory. In fact, if you created a bar graph of the assembly language instructions used in a program, whether the original program was written in C++ or assembly language, you'll find that the most used instruction is the MOVE instruction. Therefore, one of our first tasks is to build on our understanding of how memory systems are organized, and then look at the memory system from the point-of-view of the processor.

Figure 7.1 is a fairly typical memory map for a 68K processor. The processor can uniquely address 16,777,216 (16M) bytes of memory. It has 23 external word address lines and two byte selector control signals for choosing one or the other of the two bytes stored in a word location. The external data bus is 16-bits wide, but the internal data paths are all 32 bits wide. Subsequent members of the 68K family, beginning with the 68020, all had 32-bit external data buses. In figure 7.1 we see that the first 128K words (16-bit) of memory occupy the address range from 0x000000 to 0x01FFFE. This will typically be the program code and the vector tables. The vector table occupies the first 256 long words of memory. This is the reason that non-volatile, read-only memory occupies the lower address range. The upper address range usually contains the RAM memory. This is volatile, read/write memory where variables are stored. The rest of the memory space contains the address of I/O devices. All the rest of the memory in figure 7.1 is empty space. Also, we'll see in a moment why the addresses for the last word of ROM and the last word of RAM are 0x01FFFE and 0xFFFFFE, respectively.

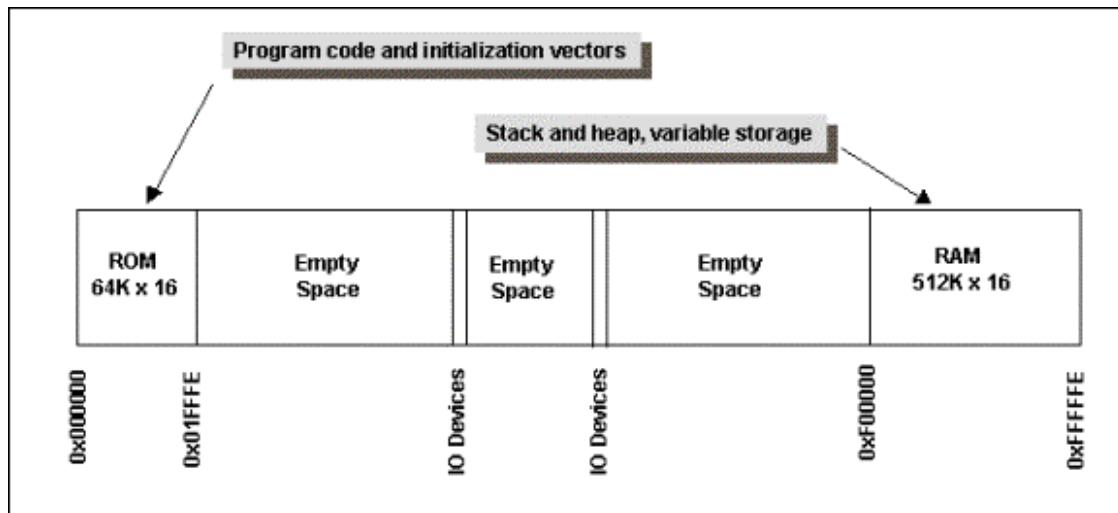


Figure 7.1: Memory map for a 68K-based computer system.

Since 8-bits, or a byte, is the smallest data quantity that we usually deal with, storing byte-sized characters (chars) in memory is straightforward when we're dealing with a memory that is also 8-bits wide. However, our PCs all have memories that are 32 bits wide and lots of computers are in use today that have 16 bit data path widths. As an extreme example, we don't want to waste 24 bits of storage space simply because we want to store an 8-bit quantity in a 32-bit memory, so computers are designed to allow **byte addressing**. Think of byte addressing in a 32-bit word as an example of paging with the page address being the 32-bit word and the offset being the 4 possible byte positions in the page. However, there is one major difference between byte addressing and real paging. With byte addressing, we do not have a separate word address which corresponds to a page address. Figure 7.2 shows this important distinction.

We call this type of memory storage **byte packing** because we are literally packing the 32-bit memory full of bytes. This type of addressing introduces several ambiguities. One is quite serious, and the others are merely new to us. We discuss the serious one in a moment. Referring to figure 7.2 we see that the byte at memory address FFFF0 (char) and the 32-bit long word (int) at FFFF0 have the same address. Isn't this a disaster waiting to happen? The answer is a definite "maybe". In C and C++, for example, you must declare a variable and its type before you can use it. Now you see the reason for it. Unless the compiler knows the type of the variable stored at address FFFF0, it doesn't know the type of code it must generate to manipulate it. Also, if you want to store a char at FFFF0, then the compiler has to know how much storage space to allocate to it.

Also, notice that we cannot access 32-bit words at addresses other than those that are divisible by 4. Some processors will allow us to store a 32-bit value at an odd boundary, such as byte addresses 000003-000006, but many others will not. The reason is that the processor would have to make several additional memory operations to read in the entire value and then do some extra work to reconstruct the bytes into the correct order. We call this a **non-aligned access** and it is generally quite costly in terms of processor performance to allow it to happen. In fact, if you look at the memory map of how a

compiler stores objects and structures in memory, you'll often see spaces in the data, corresponding to intentional gaps so as not to create regions of data containing non-aligned accesses.

<u>Word Address</u>				
000000	Byte 0 – Address 000000	Byte 1 – Address 000001	Byte 2 – Address 000002	Byte 3 – Address 000003
000004	Byte 0 – Address 000004	Byte 1 – Address 000005	Byte 2 – Address 000006	Byte 3 – Address 000007
000008	Byte 0 – Address 000008	Byte 1 – Address 000009	Byte 2 – Address 00000A	Byte 3 – Address 00000B
00000C	Byte 0 – Address 00000C	Byte 1 – Address 00000D	Byte 2 – Address 00000E	Byte 3 – Address 00000F
000010	Byte 0 – Address 000010	Byte 1 – Address 000011	Byte 2 – Address 000012	Byte 3 – Address 000013
● ● ●				
FFFFF0	Byte 0 – Address FFFF00	Byte 1 – Address FFFF01	Byte 2 – Address FFFF02	Byte 3 – Address FFFF03
FFFFF4	Byte 0 – Address FFFF40	Byte 1 – Address FFFF41	Byte 2 – Address FFFF42	Byte 3 – Address FFFF43
FFFFF8	Byte 0 – Address FFFF80	Byte 1 – Address FFFF81	Byte 2 – Address FFFF82	Byte 3 – Address FFFF83
FFFFFC	Byte 0 – Address FFFF0C	Byte 1 – Address FFFF0D	Byte 2 – Address FFFF0E	Byte 3 – Address FFFF0F

Figure 7.2: Relationship between word addressing and byte addressing in a 32-bit wide word.

Also notice that when we are doing 32-bit word accesses, address bits A0 and A1 aren't being used. This might prompt you to ask, "If we don't use them, what good are they?" However, we do need them when we need to access a particular byte within the 32-bit words. A0 and A1 are often called the **byte selector** address lines because that is their main function. Another point is that we really only need byte selectors when we are writing to memory. Reading from memory is fairly harmless, but writing changes everything. Therefore, you want to be sure that you modify only the byte you are interested in and not the others. From a hardware designer's perspective having byte selectors allows you to qualify the write operation to only the byte that you are interested in.

Many processors will not explicitly have the byte selector address lines at all. Rather, they provide signals on the status bus which are used to qualify the WRITE operations to memory. What about storing 16-bit quantities (a **short** data type) in 32-bit memory locations? The same rules apply in this case. The only valid addresses would be those addresses divisible by 2, such as 000000, 000002, 000004, etc. In the case of 16-bit word addressing, the lowest order address bit, A0, isn't needed. For our 68K processor, which has a 16-bit wide data bus to memory, we can store two bytes in each word of memory, so A0 isn't used for word addressing and becomes the byte selector for the processor.

Figure 7.3 shows a typical 32-bit processor and memory system interface. The READ signal from the processor and the CHIP SELECT signals have been omitted for clarity.

The processor has a 32-bit data bus and a 32-bit address bus. The memory chips represent one page of RAM somewhere in the address space of the processor. The exact page of memory would be determined by the design of the Address Decoder logic block. The RAM chips each have a capacity of 1 M bit and are organized as 128K by 8.

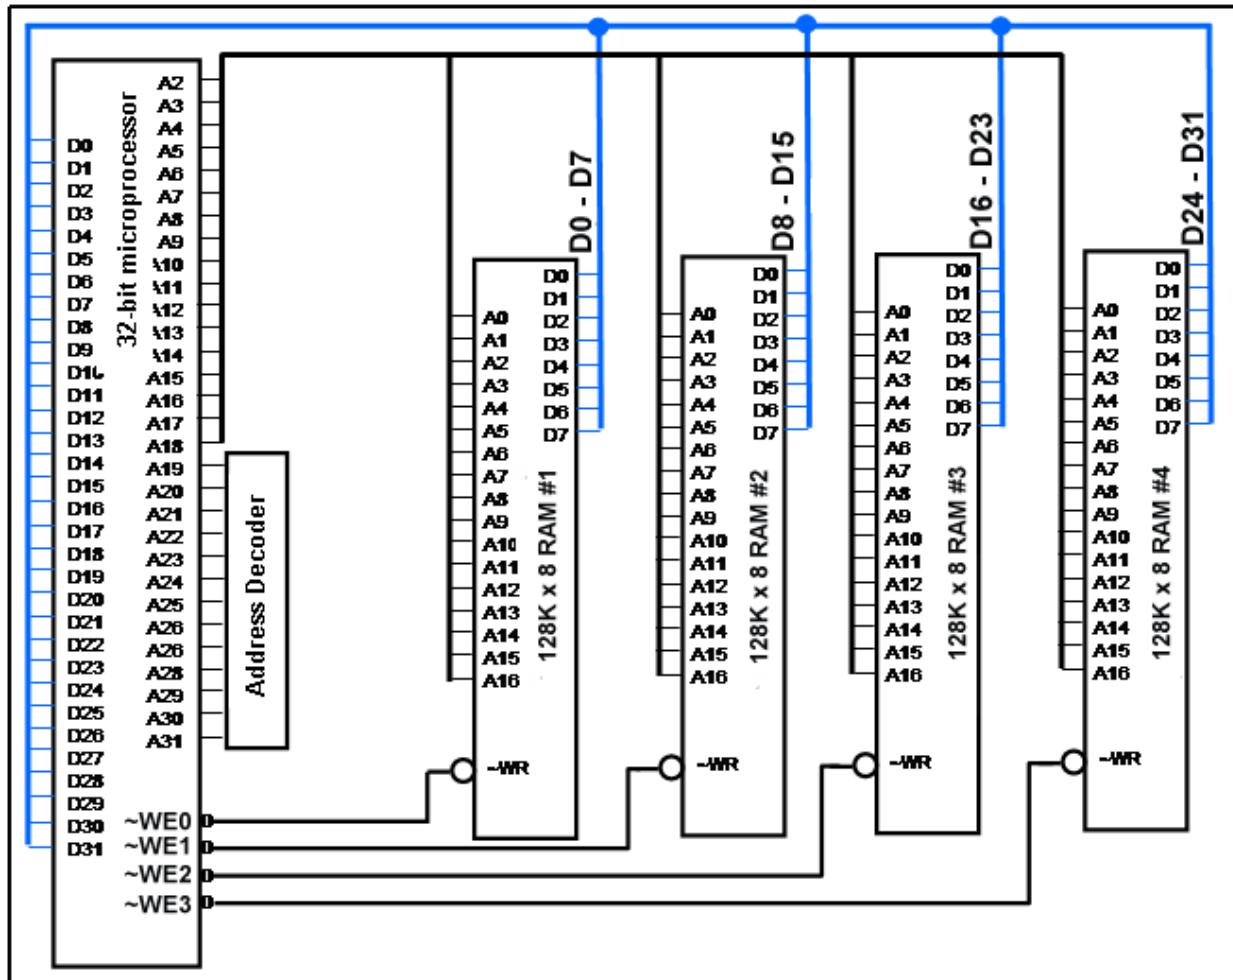


Figure 7.3: Memory organization for a 32-bit microprocessor. CHIP SELECT and READ signals have been omitted for clarity. Note that address lines A0 and A1 do not exist in this diagram because they are synthesized by the four WRITE ENABLE lines, $\sim\text{WE}0$ through $\sim\text{WE}3$.

Since we have a 32-bit wide data bus and each RAM chip has 8 data I/O lines, we need 4 memory chips per 128K wide page. Chip #1 is connected to data lines D0 through D7, chip #2 is connected to data lines D8 through D15, chip #3 is connected to data lines D16 through D23 and chip #4 is connected to data lines D24 to D31, respectively.

The address bus from the processor contains 30 address lines, which means it is capable of addressing 2^{30} long words (32-bit wide). The additional addressing bits needed to address the full address space of 2^{32} bytes are implicitly controlled by the processor internally and explicitly controlled through the 4 WRITE ENABLE signals labeled $\sim\text{WE}0$ through $\sim\text{WE}3$.

Address lines A2 through A18 from the processor are connected to address inputs A0 through A16 of the RAM chips, with A2 from the processor being connected to A0 on each of the 4 chips, and so on. This may seem odd at first, but it should make sense to you after you think about it. In fact, there is no special reason that each address line from the processor must be connected to the same address input pin on each of the memory devices. For example, A2 from the processor could be connected to A14 of chip #1, A3 of chip #2, A8 of chip #3 and A16 of chip #4. The same address from the processor would clearly be addressing different byte addresses in each of the 4 memory chips, but as long as all of the 17 address lines from the processor are connected to all 17 address lines of the memory devices, the memory should work properly.

The upper address bits from the processor, A19 through A31 are used for the page selection process. These signals are routed to the address decoding logic where the appropriate ~CHIP SELECT signals are generated. These signals have been omitted from figure 7.3. There are 13 higher order address bits in this example. This gives us 2^{13} or 8,192 pages of memory. Each page of memory holds 128K 32-bit wide words, which works out to be 2^{30} long words. In terms of addresses, each page actually contains 512K bytes, so the byte address range on each page goes from byte address (in HEX) 00000 through 7FFFF. Recall that in this memory scheme, there are 8,192 pages with each page holding 512K bytes. Thus, page addresses go from 0000 through 1FFF. It may not seem obvious to you to see how the page addresses and the offset addresses are related to each other but if you expand the hexadecimal addresses to binary and lay them next to each other you should see the full 32-bit address.

Now, let's see how the processor deals with data sizes smaller than long words. Assume that the processor wants to read a long word from address ABCDEF64 and that this address is decoded to be on the page of figure 7.3. Since this address is on a 32 bit boundary, A0 and A1 = 0, and are not used as part of the external address that goes to memory. However, if the processor wanted to do a word access of either one of the words located at address ABCDEF64 or ABCDEF66, it would still generate the same external address. When the data was read into the processor, the $\frac{1}{2}$ of the long word that was not needed would be discarded. Since this is a READ operation, the contents of memory are not affected.

If the processor wanted to read any one of the 4 bytes located at byte address ABCDEF64, ABCDEF65, ABCDEF66 or ABCDEF67, it would still perform the same read operation as before. Again, only the byte of interest would be retained and the others would be discarded.

Now, let's consider a write operation. In this case, we are concerned about possibly corrupting memory, so we want to be sure that when we write a quantity smaller than a long word to memory, we do not accidentally write more than we intend to. So, suppose that we only want to write a new value to the byte at memory location ABCDEF65. In this situation, the ~WE1 signal would have to be asserted, so only the byte at byte address 1 position could be modified.

Thus, to write a byte to memory, we only activate one of the 4 WRITE ENABLE signals.

To write a word to memory we would activate either ~WE0 and ~WE1 together or ~WE2 and ~WE3 together. Finally, to write a long word, all 4 of the WRITE ENABLE lines would be asserted.

What about the case of a 32-bit word stored in a 16-bit memory? In this case, the 32-bit word can be stored on any even word boundary because the processor must always do two consecutive memory accesses to retrieve the entire 32-bit quantity. However, most compilers will still try to store the 32-bit words on natural boundaries (addresses divisible by 4). This is why assembly language programmers can often save a little space or speed up an algorithm by overriding what the compiler does to generate code and tweaking it for greater efficiency.

Let's get back on track. For a 32-bit processor, address bits A2 . . . A31 are used to address the 1,073,741,824 possible long words, and A0 . . . A1 address the four possible bytes within the long word. This gives us a total of 4,294,967,296 addressable byte locations in a 32-bit processor. In other words, we have a byte addressing space of 4 GB. A processor with a 16-bit wide data bus, such as the 68K, uses address lines A1-A23 for word addressing and A0 for byte selection.

Combining all of this, you should see the problem. You could have an 8-bit byte, a 16-bit word or a 32-bit word with the same address. Isn't this ambiguous? Yes it is. When we're programming in a high-level language, we depend upon the compiler to keep track of these messy details. This is one reason why *casting* one variable type to another can be so dangerous. When we are programming in a low-level language, we depend upon the skill of the programmer to keep track of this.

Seems easy enough, but it's not. This little fact of computer life is one of the major causes of software bugs. How can a simple concept be so complex? It's not complex, it's just ambiguous. Figure 7.4 illustrates the problem. The leftmost column of figure 7.4 shows a string (aptly named "string") stored in an 8-bit memory space. Each ASCII character occupies successive memory locations.

The middle column shows a 16-bit memory that is organized so that successive bytes are stored right to left. The byte corresponding to A0 = 0 is aligned with the low order portion, DB0 . . . DB7, of the 16-bit word and the byte corresponding to A0 = 1 is aligned with the high order portion, DB8..DB15, of the 16-bit word. This is called **Little Endian** organization. The rightmost column stores the characters as successive bytes in a left to right fashion. The byte position corresponding to A0 = 0 is aligned with the high order portion of the 16-bit word. This is called **Big Endian** organization. As an exercise, consider how the bytes are stored in figure 7.2. Are they Big or Little Endian?

Motorola and Intel chose to use different endian conventions and Pandora's Box was opened for the programming world. Thus, C or C++ code written for one convention would have subtle bugs when ported to the other convention. It gets worse than that. Engineers working together on projects misinterpret specifications if the intent is one convention and they assume the other. The ARM architecture allows the programmer to

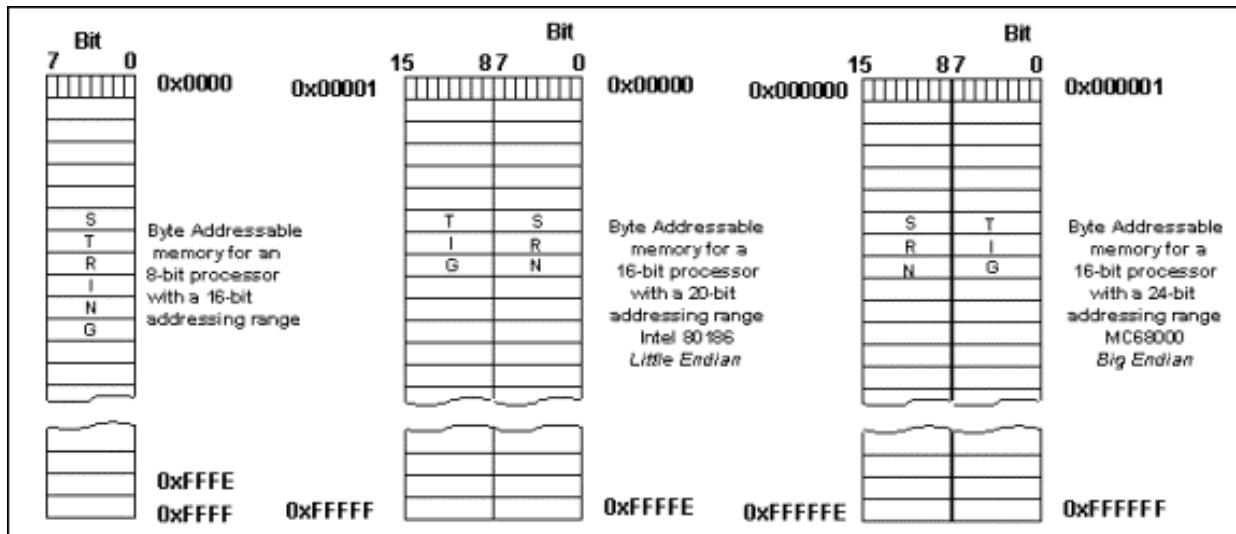


Figure 7.4: Two methods of packing bytes into 16-bit memory words. Placing the low order byte at the low order end of the word is called *Little Endian*. Placing the low order byte at the high order side of the word is called *Big Endian*.

establish which type of “endianess” will be used by the processor at power-up. Thus, while the ARM processor can deal with either big or little endian, it cannot dynamically switch modes once the endianess is established. Figure 7.5 shows the difference between the two conventions for a 32-bit word packed with four bytes.

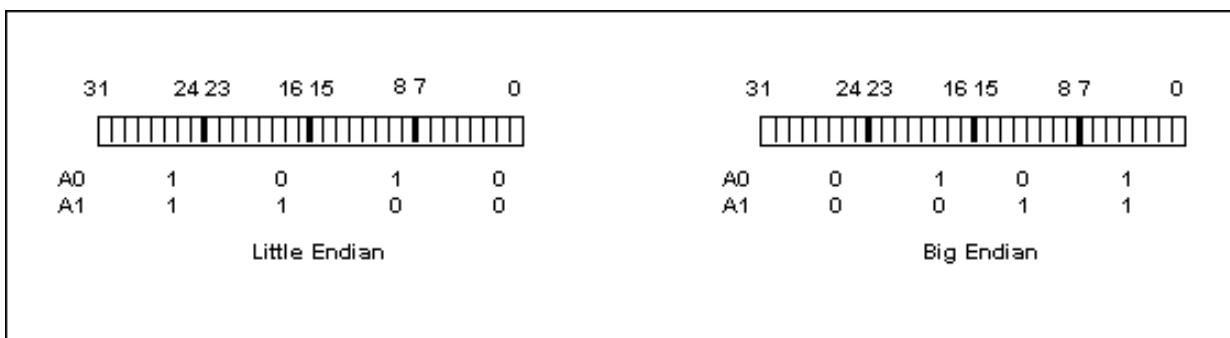


Figure 7.5: Byte packing a 32-bit word in Little Endian and Big Endian modes.

If you take away anything from this text, remember this problem because you will see it at least once in your career as a software developer. Before you accuse me of beating this subject to death, let's look at it one more time from the hardware perspective. The whole

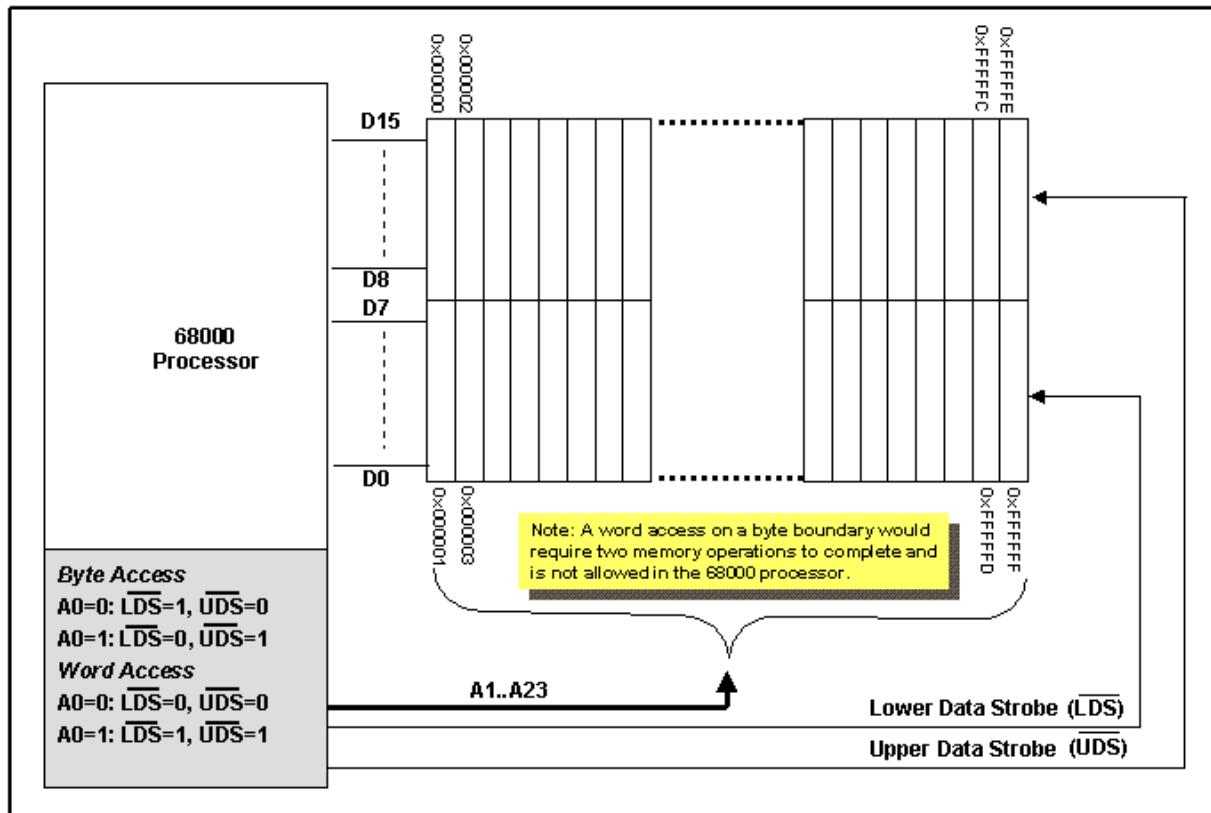


Figure 7.6: Memory addressing modes for the Motorola 68K processor

area of memory addressing can be very confusing for novice programmers as well as seasoned veterans. Also, there can be ambiguities introduced by architectures and manufacturer's terminology. So, let's look at how Motorola handled it for the 68K and perhaps this will help us to better understand what's really going on, at least in the case of the Motorola processor, even though we have already looked at the problem once before in figure 7.3. Figure 7.6 summarizes the memory addressing scheme for the 68K processor.

The 68K processor is capable of directly addressing 16 Mbytes of memory, requiring 24 “effective” addressing lines. Why? Because $2^{24} = 16,777,216$. In figure 7.6 we see 23 address lines. The missing address line, A0, is synthesized by two additional control signals, ~LDS and ~UDS.

For a 16-bit wide external data bus, we would normally address bit A0 to be the byte selector. When A0 is 0, we choose the even byte, and when A0 = 1, we choose the odd byte. The *endianness* of the 68K is Big Endian, so that the even byte is aligned with D8

through D15 of the data bus. Referring to figure 7.6, we see that there are two status bus signals coming out of the processor, designated $\sim\text{UDS}$, or *Upper Data Strobe*, and $\sim\text{LDS}$, or *Lower Data Strobe*.

When the processor is doing a byte access to memory, then either $\sim\text{LDS}$ or $\sim\text{UDS}$ is asserted to indicate to the memory which part of the word is being accessed. If the byte at the even address is being accessed ($A_0 = 0$), then UDS is asserted and LDS stays HIGH, or OFF state. If the odd byte is being accessed ($A_0 = 1$), then $\sim\text{LDS}$ is asserted and $\sim\text{UDS}$ remains HIGH. For a word access, both $\sim\text{UDS}$ and $\sim\text{LDS}$ are asserted. This behavior is summarized in the table of Figure 7.6.

You would normally use $\sim\text{LDS}$ and $\sim\text{UDS}$ as gating signals to the memory control system. For example, you could use the circuit shown in figure 7.7 to control which of the bytes are being written to.

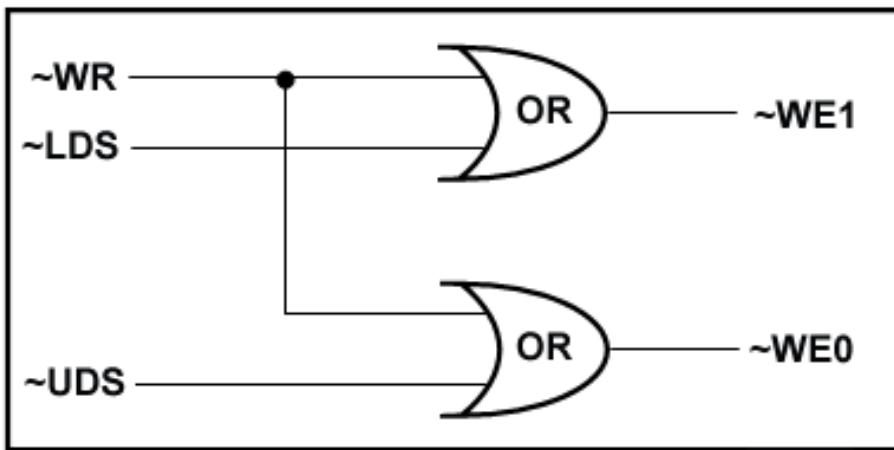


Figure 7.7: Simple circuit to control the byte writing of a 68K processor

You may be scratching your head about this circuit. Why are we using OR gates? We can answer the question in two ways. First, since all of the signals are asserted LOW, we are really dealing with the negative logic equivalent of an AND function. The gate that happens to be the negative logic equivalent of the AND gate is the OR gate, since the output is 0 if and only if both inputs are 0.

The second way of looking at it is through the equivalence equations of DeMorgan's Theorems. Recall that:

$$\sim(A * B) = \sim A + \sim B \quad (1)$$

$$\sim(A + B) = \sim A * \sim B \quad (2)$$

In this case, equation 1 shows that the OR of $\sim A$ and $\sim B$ would be equivalent to using positive logic A and B and then obtaining the NAND of the two signals.

Now, suppose that you attempted to do a word access at an odd address. For example, suppose that you wrote the following assembly language instruction:

```
move.w D0,$1001
```

*** This is a non-aligned access!**

This instruction tells the processor to make a copy of the word stored in internal register D0 and store the copy of that word in external memory beginning at memory address \$1001. The processor would need to execute two memory cycles to complete the access because the byte order requires that it bridge the two memory locations to correctly place the bytes. Some processors are capable of this type of access, but the 68K is one of the processors that can't. If a non-aligned access occurs, the processor will generate an exception and automatically branch to user-defined code (hopefully, the programmer has provided for this situation) that will attempt to deal with the error, or at least die gracefully. Since the 68K processor is a 32-bit wide processor internally, but with only a 16-bit wide data bus to external memory, we need to know how it stores 32-bit quantities in external memory as well. Figure 7.8 shows us the convention for storing long words in memory.

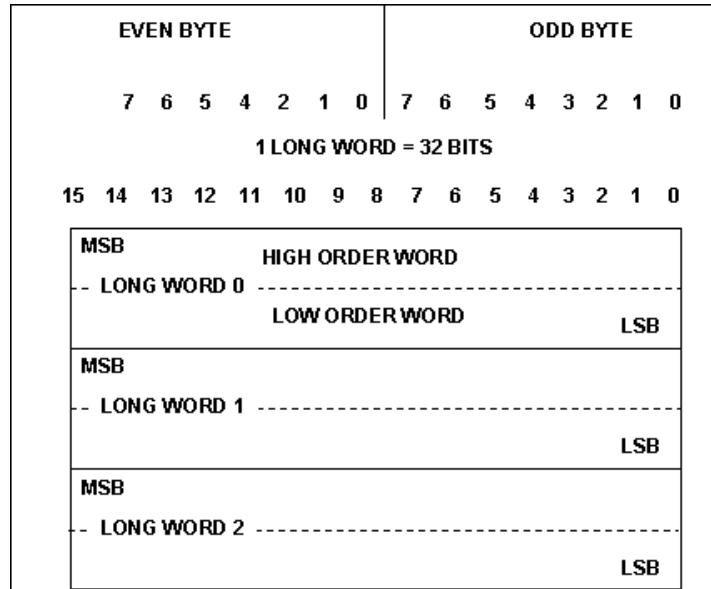


Figure 7.8: Memory storage conventions for the Motorola 68000 processor

Although figure 7.8 may seem confusing, it is just a restatement of figure 7.2 in a somewhat more abbreviated format. The figure tells us that 32-bit data quantities, called **longs** or **long words**, are stored in memory with the most significant 16 bits (D16 - D31) stored in the first word address location and the least significant 16 bits (D0 - D15) stored in the next highest word location. Also, even byte addresses are aligned with the high order portion of the 16-bit word address (Big Endian).

Introduction to Assembly Language

The PC world is dominated by an instruction set architecture (ISA) first defined by Intel over twenty-five years ago. This architecture, called X86 because of its family members, has the following lineage:

8080 → 8086 → 80186 → 80286 → 80386 → 80486 → Pentium

The world of embedded microprocessors—that is microprocessors used for a single purpose inside a device, such as a cell phone—is dominated by the Motorola 680X0 ISA:

68000 → 68010 → 68020 → 68030 → 68040 → 68060 → ColdFIRE

The ColdFIRE instruction set architecture is a modern microprocessor architecture, called RISC. We'll discuss the RISC architecture in a later chapter. Yet, it maintains backward compatibility with the original 68K ISA. Backward compatibility is very important because there is so much 68K code still around and being used. The Motorola 68K instruction set is one of the most studied ISAs around and you can find an incredible number of hits if you do a Web search on "68K" or "68000."

Every computer system has a fundamental set of operations that it can perform. These operations are defined by the **instruction set** of the processor. The reason for a particular set of instructions is due to the way the computer is organized internally and designed to operate. This is what we would call the architecture of the computer. The architecture is mirrored by the assembly language instructions that it can execute, because these

Instead of writing a program in machine language as:	We write the program in assembly language as:
<pre>00000412 307B704B 00000416 327B704A 0000041A 1080 0000041C B010 000041E 67000008 00000422 1600 00000424 61000066 00000428 5248 0000042A B0C9</pre>	<pre>MOVEA.W (TEST_S,PC,D7),A0 "We'll use address indirect MOVEA.W (TEST_E,PC,D7),A1 "Get the end address MOVE.B D0,(A0) "Write the byte CMP.B (A0),D0 "Test it BEO NEXT_LOCATION "OK, keep going MOVE.B D0,D3 "copy bad data BSR ERROR "Bad byte ADDQ.W #01,A0 "increment the address CMPA.W A1,A0 "are we done?</pre>

Figure 7.9: The box on the right is a snippet of 68K code in assembly language. The box on the left is the machine language equivalent.

instructions are the mechanism by which we access the computer's resources. The instruction set is the atomic element of the processor. All of the complex operations are achieved by building sequences of these fundamental operations.

Computers don't read assembly language. They take their instructions in machine code. As you know, the machine code defines the entry point into the state machine microcode table that starts the instruction execution process. Assembly language is the **human-readable** form of these machine language instructions. There is nothing mystical about the machine language instructions, and pretty soon you'll be able to understand them and see the patterns that guide the internal state machines of the modern microprocessor. For now, we'll focus on the task of learning assembly language. Consider figure 7.9.

The machine language code is actually the output of the assembler program that converts the assembly language source file, which you write with a text editor, into the

corresponding machine language code that can be executed by a 680X0 processor. The left hand box actually has two columns, although it may be difficult to see that. The left column starts with the hexadecimal memory location where the machine language instruction is stored. In this case the memory location \$00000412 holds the machine language instruction code 0x307B7048. The next instruction begins at memory location 0x00000416 and contains the instruction code 0x327B704A. These two machine language instructions are given by these assembly language instructions.

```
MOVEA.W (TEST_S, PC, D7), A0
```

```
MOVEA.W (TEST_E, PC, D7), A1
```

Soon you'll see what these instructions actually mean. For now, we can summarize the above discussion this way:

- Starting at memory location \$00000412¹, and running through location \$00000415, is the machine instruction code \$307B7048. The assembly language instruction that corresponds to this machine language data is **MOVEA.W (TEST_S, PC, D7), A0**
- Starting at memory location \$00000416, and running through location \$00000419, is the machine instruction code \$327B704A. The assembly language instruction that corresponds to this machine language data is **MOVEA.W (TEST_E, PC, D7), A1**

Also, for the 68K instruction set, the smallest machine language instruction is 16-bits long (4 hex digits). No instruction will be smaller than 16-bits long, although some instructions may be as long as 5, 16-bit words long.

There is a 1:1 correspondence between assembly language instructions and machine language instructions. The assembly language instructions are called ***mnemonics***. They are designed to be a shorthand clue as to what the instruction actually does. For example:

MOVE.B	move a byte of data
MOVEA.W	move a word of data to an address register
CMP.B	compare the magnitude of two bytes of data
BEQ	branch to a different instruction if the result equals zero
ADDQ.W	add (quickly) two values
BRA	always branch to a new location

You'll notice that I've chosen a different font for the assembly language instructions. This is because fonts with fixed spacing, like "Courier", keep the characters in column

¹ In standardized languages like C or C++ hexadecimal numbers are represented with the prefix "0x". However, there is no such standard for assembly language. We will adopt the Motorola standard prefix which uses the dollar sign, "\$", to indicate a hexadecimal number.

alignment, which makes it easier to read assembly language instructions. There's no law that you must use this font, the assembler probably doesn't care, but it might make it easier for you to read and understand your programs if you do.

The part of the instruction that tells the computer what to do is called the *opcode* (short for "operation code"). This is only one half of the instruction. The other half tells the computer how and where this operation should be performed. The actual opcode, for example MOVE.B, is actually an opcode and a modifier. The opcode is MOVE. It says that some data should be moved from one place to another. The modifier is the ".B" suffix. This tells it to move a byte of data, rather than a word or long word. In order to complete the instruction we must tell the processor:

1. where to find the data (this is called operand 1), and
2. where to put the result (operand 2).

A complete assembly language instruction must have an opcode and may have 0,1 or 2 operands. Here's your first assembly language instruction. It's called NOP (pronounced *No op*). It means do nothing. You might be questioning the sanity of this instruction but it is actually quite useful; compilers make very good use of them. The NOP instruction is an example of an instruction that takes 0 operands.

The instruction CLR.L D4 is an example of an instruction that takes one operand. It means to clear, or set to zero, all 32 bits (the ".L" modifier) of the internal register, D4.

The instruction MOVE.W D0,D3 is an example of an instruction that takes two operands. Note the comma separating the two operands, D0 and D3. The instruction tells the processor to move 16 bits of data (the ".W" modifier) from data register D0 to data register D3. The contents of D0 are not changed by the operation. All assembly language programs conform to the following structure:

Column 1	Column 2	Column 3	Column 4 . . .
LABEL	OPCODE	OPERAND1, OPERAND2	*COMMENT

Each instruction occupies one line of text, starting in column 1 and going up to 132 columns.

1. The LABEL field is optional, but it must always start in the first column of a line. We'll soon see how to use labels.
2. The OPCODE is next. It must be separated from the label by white space, such as a TAB character or several spaces, and it must start in column 2 or later.
3. Next, the operands are separated from the opcode by white space, usually a tab character. The two operands should be separated from each other by a comma. There is no white space between the operands and the comma.

- The comment is the last field of the line. It usually starts with an asterisk or a semi-colon, depending upon which assembler is being used. You can also have comment lines, but then the asterisk must be in column 1.

Label

Although the label is optional, it is a very important part of assembly language programming. You already know how to use labels when you give a symbolic name to a variable or a constant. You also use labels to name functions. In assembly language we

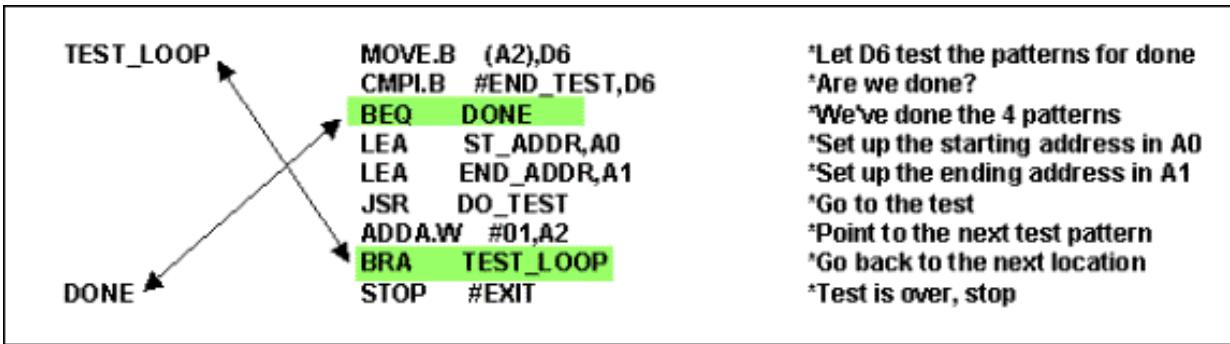


Figure 7.10: Snippet of assembly language code demonstrating the use of labels.

commonly use labels to refer to the memory address corresponding to an instruction or data in the program. The label must be defined in column 1 of your program. Labels make the program much more readable. It is possible to write a program without using labels, but almost no one ever does it that way. The label allows the assembler program to automatically (and correctly!) calculate the addresses of operands and destinations. For example, consider the following snippet of code in figure 7.10.

The code example of figure 7.10 has two labels, TEST_LOOP and DONE. These labels correspond to the memory locations of the instructions, "MOVE.B (A2),D6" and "STOP #EXIT" respectively. As the assembler program converts the assembly language instructions into machine language instructions it keeps track of where in memory each instruction will be located. When it encounters a label as an operand, it replaces the label text with the numeric value. Thus, the instruction "BEQ DONE" tells the assembler to calculate the numeric value necessary to cause the program to jump to the instruction at the memory location corresponding to the label "DONE" if the test condition, equality, is met. We'll soon see how to test this equality. If the test fails, the branch instruction is ignored and the next instruction is executed.

Comments

Before we get too far off-shore, we need to make a few comments about the proper form for commenting your assembly language program. As you can see from figure 7.10 each assembly language instruction has a comment associated with it. Different assemblers handle comments in different ways. Some assemblers require that comments that are on a line by themselves have an asterisk '*' or a semicolon ';' as the first character on the line. Comments that are associated with instructions or assembler directives might need the

semicolon or asterisk to begin the comment, or they might not need any special character because the preceding white space defines the location of the comment block. The important point is that assembly language code is not self-documenting, and it is easy for you to forget, after a day or so goes by, exactly what you were trying to do with that algorithm.

Assembly code should be profusely commented, not only for your sanity, but for the people who will have to maintain your code after you move on. There is no reason that assembly code cannot be as easy to read as a well-documented C++ program. Use equates and labels to eliminate magic numbers and to help explain what the code is doing. Use comment blocks to explain what you are doing in each section of an algorithm and what assumptions you are making about the algorithm. Finally, comment each instruction, or small group of instructions, in order to make it absolutely clear what is going on.

In his book, ***Hackers: Heroes of the Computer Revolution***, Steven Levy¹ describes the coding style of Peter Samson, an MIT student, and early programmer,

*...Samson, though, was particularly obscure in refusing to add comments to his source code, explaining what he was doing at a given time. One well-distributed program Samson wrote went on for several hundreds of assembly language instructions, with only one comment beside an instruction which contained the number 1750. The comment was **RIPJSB**, and people racked their brains about its meaning until someone figured out that 1750 was the year that Bach died, and that Samson had written an abbreviation for Rest In Peace Johann Sebastian Bach.*

Programmer's Model Architecture

In order to program in assembly language, we must be familiar with the basic architecture of the processor. Our view of the architecture is called the ***Programmer's Model*** of the processor. We must understand two aspects of the architecture:

1. The instruction set, and
2. The ***addressing modes***.

The addressing modes of a computer describe the different ways in which it accesses the operands, or retrieves the data to be operated on. Then, the addressing modes describe what to do with the data after the operation is completed. The addressing modes also tell the processor how to calculate the destination of a non-sequential instruction fetch, such as a branch or jump to a new location. Addressing modes are so important to the understanding of the computer we'll need to study them a bit before we can create an assembly language program. Thus, unlike C, C++, or JAVA, we'll need to develop a

certain level of understanding for the machine that we're programming before we can actually write a program.

Unlike C or C++, assembly language is not portable between computers. An assembly language program written for an Intel 80486 will not run on a Motorola 68000. A C program written to run on an Intel 80486 *might be able to run* on a Motorola 68000 once the original source code is recompiled for the 68000 instruction set, but differences in the architectures, such as big and little endian, may cause errors to be introduced.

It might be worthwhile and stop for a moment in order to reflect on why, as a programmer, it is important to learn assembly language. Computer science and programming depends upon a working knowledge of the processor, its limitations and its strengths. An understanding of assembly language provides the critical insight into the computing engine that is executing your code. Even though high-level languages like C++ and JAVA do a good job of abstracting the low level details, it is important to keep in mind that the engine is not infinitely powerful and that its resources are not limitless.

Assembly language is tightly coupled to the design of the processor and represents the first level of simplification of the binary instruction set architecture into a human readable form. Generally there is a 1:1 match between the assembly language instruction and the binary or hexadecimal instruction that results. This is very different from C or C++, where one C statement, may generate hundreds of lines of assembly code.

It is still true that you can generally write the tightest code in assembly language. While C compilers have gotten pretty good, they're not perfect. A program written in assembly language will often use less space and run faster than a program with the same functionality written in C. Many games are still coded in assembly language for exactly that reason.

The need for efficient code is especially true of interrupt handlers and algorithms written for specialized processors, such as *digital signal processors (DSPs)*. Many experienced programmers will argue that any code that must be absolutely deterministic cannot be written in C, because you cannot predict the execution time for the code generated by the compiler. With assembly language you can control your program at the level of a single clock cycle. Also, certain parts of the start-up code for the C runtime environment must be written in assembly because programs written in C or C++ can't run properly unless the run-time environment is established before the C code can begin to execute. Thus, boot-up code tends to be written in assembly. Finally, understanding assembly language is critically important for debugging real time systems. If you've ever programmed in an environment such as Microsoft's Visual C++® and inadvertently stepped into a library function, you will then have found yourself knee-deep in x86 ISA and assembly language.

Motorola 68000 Microprocessor Architecture

Figure 7.11 is a simplified schematic diagram of the 68K architecture. Since Motorola has a large number of family members, this particular architecture is also referred to as CPU16. CPU16 is a subset of the CPU32 architecture, the ISA for the 68020 and later processors.

Let's briefly identify some of the important functional blocks that we'll later be using.

- **Program counter:** Used to hold the address of the next instruction to be fetched from memory. As soon as the current instruction is decoded by the processor, the program counter (PC) is updated to point to the address of the next sequential

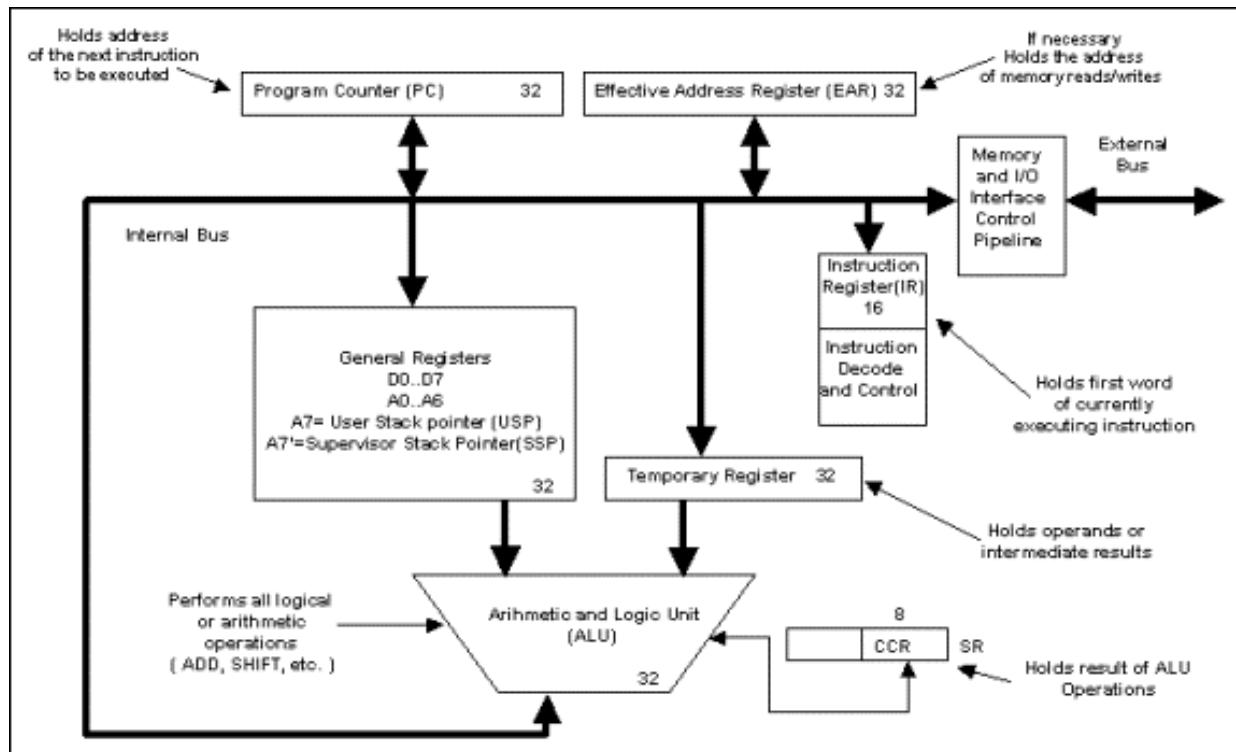


Figure 7.11: Architecture of the Motorola 68K processor.

- instruction in memory.
- **General registers:** The 68K processor has 15 general-purpose registers and two special-purpose registers. The general purpose registers are further divided into eight data registers, D0 . . . D7 and seven address registers, A0 . . . A6. The data registers are used to hold and manipulate data variables and the address registers are used to hold and manipulate memory addresses. The two special purpose registers are used to implement two separate stack pointers, A7 and A7'. We'll discuss the stack a bit later in the text.
- **Status register:** The status register is a 16-bit register. The bits are used to describe the current state of the computer on an instruction-by-instruction basis. The condition code register, CCR, is part of the status register and holds the bits that directly relate to the result of the last instruction executed.

- **Arithmetic and logic unit (ALU):** The ALU is the functional block where all of the mathematical and logical data manipulations are performed.

Figure 7.12 is the *Programmer's Model* of the 68K architecture. It differs from figure 7.11 because it focuses only on the details of the architecture that is relevant to writing a program.

In this text, we're not going to deal with the status register (SR), or the supervisor stack pointer (A7'). Our world from now on will focus on D0 . . . D7, A0 . . . A6, A7, CCR and the PC. From these registers, we'll learn everything that we need to know about this computer architecture and assembly language programming.

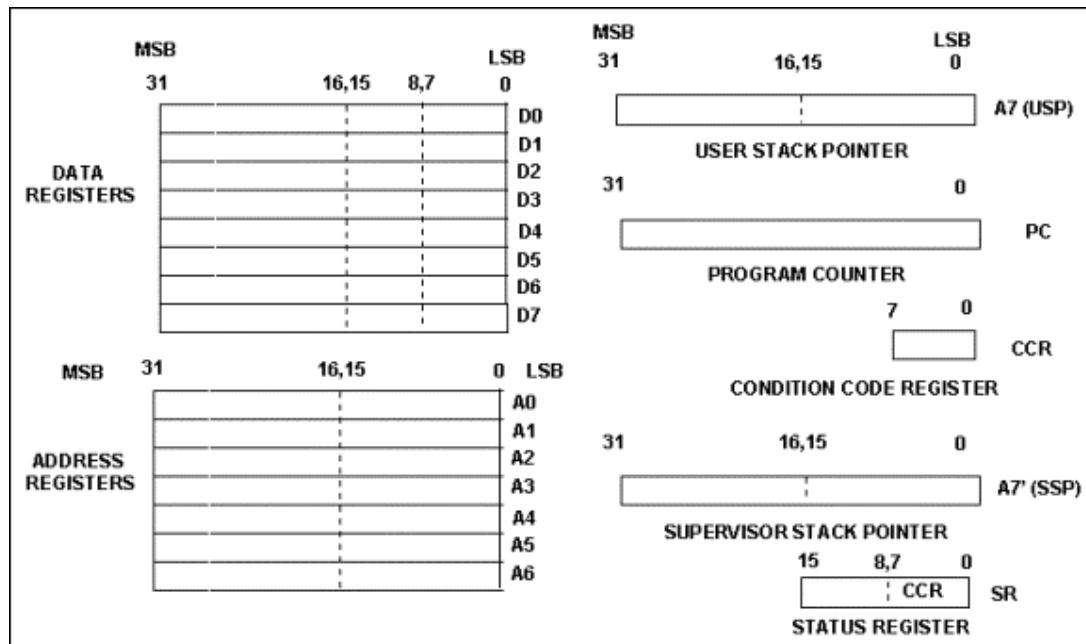


Figure 7.12: Programmer's model of the 68K processor.

Condition Code Register

The condition code register, CCR, deserves some additional explanation. The register is shown in more detail in figure 7.13.

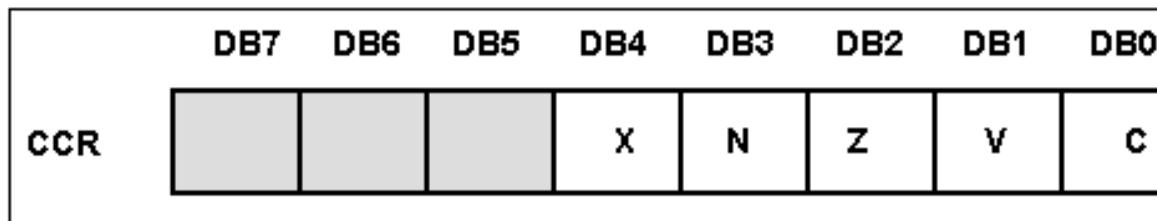


Figure 7.13: Condition Code Register. The shaded bits are not used.

The CCR contains a set of five condition bits whose value can change with the result of each instruction being executed. The exact definition of each of the bits is summarized below.

- **X BIT (extend bit):** used with multi-precision arithmetic
- **N BIT (negative bit):** indicates that the result is a negative number
- **Z BIT (zero bit):** indicates that the result is equal to zero
- **V BIT (overflow):** indicates that the result may have exceeded the range of the operand
- **C BIT (carry bit):** indicates that a carry was generated in a mathematical operation

Note:

Even though the 68K processor has only 24 address lines external to the processor, internally it is still a 32-bit processor. Thus, we can represent addresses as 32-bit values in our examples.

The importance of these bits resides with the family of test and branch instructions such as BEQ, BNE, BPL, BMI, and so on. These instructions test the condition of an individual *flag*, or CCR bit, and either take the branch if the condition is true, or skip the branch and go to the next instruction if the condition is false. For example, BEQ means *branch equal*. Well, equal to what? The BEQ instruction is actually testing the state of the *zero flag*, Z. If Z = 1, it means that the result in the register is zero, so we should take the branch because the result is equal to zero. So **branch equal** is an instruction which tells the processor to take the branch if there was a recent operation that resulted in a zero result.

But BEQ means something else. Suppose we want to know if two variables are equal to each other. How could we test for equality? Simple, just subtract one from the other. If we get a result of zero (Z = 1), they're equal. If they're not equal to each other we'll get a non-zero result and (Z = 0). Thus, BEQ is true if Z = 1 and BNE (branch not equal) is true if Z = 0. BPL (branch plus) is true if N = 0 and BMI (branch minus) is true if (N = 1).

There are a total of 14 conditional branch instructions in the 68K instruction set. Some test only the condition of a single flag, others test logical combinations of flags. For example, the BLT instruction (branch less than) is defined by: $BLT = N^*V + /N^*V (N \oplus V)$ (not by bacon, lettuce, and tomato).

Effective Addresses

Let's now go back and look at the format of the instructions. Perhaps the most commonly used instruction is the MOVE instruction. You'll find that most of what you do in assembly language is moving data around. The MOVE instruction takes two operands. It looks like this:

MOVE.W source(EA),destination(EA)

For example

MOVE.W \$4000AA00,\$10003000

tells the processor to **copy** the 16-bit value stored in memory location \$4000AA00 to memory location \$10003000.

Also, the MOVE mnemonic is a bit misleading. What the instruction does is to overwrite the contents of the destination operand with the source operand. The source operand isn't changed by the operation. Thus, after the instruction both memory locations contain the same data as \$4000AA00 did before the instruction was executed.

In the above example, the source operand and the destination operand are addresses in memory that are exactly specified. These are ***absolute addresses***. They are absolute because the instruction specifies exactly where to retrieve the data from and where to place it. Absolute addressing is just one of the possible addressing modes of the 68K. We call these addressing modes ***effective addresses***. Thus, when we write the general form of the instruction:

MOVE.W source(EA),destination(EA)

we are saying that the source and destination of the data to move will be independently determined by the effective addressing mode for each.

For example:

MOVE.W D0,\$10003000

moves the contents of one of the processor's eight internal data registers, in this case data register D0, to memory location \$10003000. In this example, the mode used for the source effective address is called ***data register direct*** and the mode used for the destination effective address mode is called ***absolute***. We could also write the MOVE instruction as:

MOVE.W A0,D5

which would move the 16-bit word, currently stored in address register A0 to data register D5. The source effective address mode is ***address register direct*** and the destination effective address mode is ***data register direct***. Suppose that the contents of address register A0 is \$4000AA00 (we can write this in a shorthand notation as $\langle A0 \rangle = \$4000AA00$). The instruction

MOVE.W D5,(A0)

moves the contents of data register D5 to memory location \$4000AA00. This is an example of ***address register indirect*** addressing. You are probably already familiar with this addressing mode because this is a pointer in C++. The contents of the address register become the address of the memory operation. We call this an indirect addressing mode because the contents of the address register are not the data we want, but rather the memory address of the data we want. Thus, we aren't storing the data directly to the register A0, but indirectly by using the contents of A0 as a pointer to its ultimate destination, memory location \$4000AA00. We indicate that the address register is being used as a pointer to memory by the parentheses around the address register. Suppose that $\langle A1 \rangle = \$10003000$ and $\langle A6 \rangle = \$4000AA00$. The instruction

MOVE.W (A1),(A6)

would copy the data located in memory location \$10003000 to memory location \$4000AA00. Both source and destination effective address modes are the address register indirect mode.

Let's look at one more example.

MOVE.W #\$234A,D2

would place the hexadecimal number, \$234A directly into register D2. This is an example of the ***immediate address mode***. Immediate addressing is the mechanism by which memory variables are initialized. The pound sign (#) tells the assembler that this is a number, not a memory location. Of course, only the source effective address could be an immediate address. The destination can't be a number—it must be a memory location or a register.

The effective address (EA) specifies how the operands of the instruction will be accessed. Depending upon the instruction being executed, not all of the effective addresses may be available to use. We'll discuss this more as we go. The type of effective address that will be used by the instruction to access one or more of the operands is actually specified as part of the instruction. Recall that the minimum size for an opcode word is 16-bits in length. The opcode word provides all of the information that the processor needs to execute the instructions. However, that doesn't mean that the 16-bit opcode word, by itself, contains all the information in the instruction. It *may* contain enough information to be the entire instruction, like *NOP*, but usually it only contains enough information to know what else it must retrieve, or *fetch*, from memory in order to complete the instruction. Therefore, many instructions may be longer than the opcode word portion of the instruction.

If we think about the microcode-based state machine that drives the processor, this all begins to make sense. We need the opcode word to give us our entry point into the microcode. This is what the processor does when it is decoding an instruction. In the process of executing the entire instruction, it may need to go out to memory again to fetch additional operands in order to complete the instruction. It knows that it has to do these additional memory fetches because the path through the state machine, as defined by the opcode word, predetermines it.

Consider the form of the opcode word for the MOVE instruction shown in figure 7.14

The opcode field may contain three types of information about the MOVE instruction within the field defined by the four data bits, DB15 - DB12. These possibilities are defined as follows:

- 0001 = MOVE.B

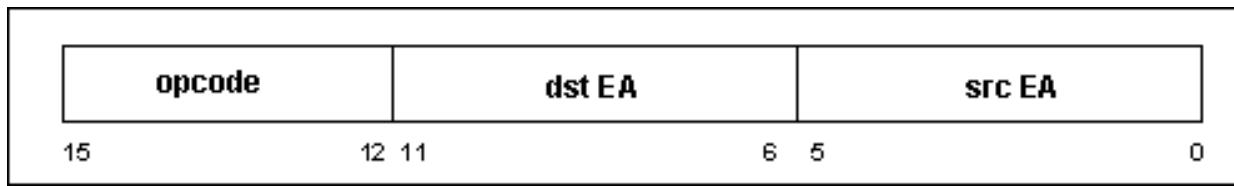


Figure 7.14: Machine language instruction format for the MOVE instruction

- 0011 = MOVE.W
- 0010 = MOVE.L

The source EA and destination EA are both 6-bit fields that are each further subdivided into two, 3-bit fields each. One 3-bit field, called the *register field*, can take on a value from 000 to 111, corresponding to one of the data registers, D0 through D7, and address register A0 through A7. The other 3-bit field is the *mode field*. This field describes which effective addressing mode is being used for the operation. We'll return to this point later on. Let's look at a real example. Consider the instruction

MOVE.W #\\$0A55,D0

What would this instruction look like after it is assembled? It may help you to refer to an assembly language programming manual, such as the *Programmer's Reference Manual*³ to better understand what's going on. Here's what we know. It's a MOVE instruction with the size of the data transfer being 16-bits, or one word. Thus, the four most significant bits of the opcode word (D15 – D12) are 0011. The source effective address (D5 – D0) is immediate, which decodes as 111 100. The destination effective address (D11 – D6) is data register D0, which decodes as 000 000. Putting this all together, the machine language opcode word is:

0011 000 000 111 100 , or \\$303C

The translation from 0011 000 000 111 100 (binary) to \$303C (hex) may have seemed a bit confusing to you because the bit fields of the instruction usually don't align themselves on nice 4-bit boundaries. However, if you proceed from right to left, grouping by four, you'll get it right every time.

Is that all the elements of the instruction? No, because all we know at this point is that the source is an immediate operand, but we still don't know what the immediate number is. Thus, the opcode word tells the processor that it has to go out to memory again and fetch another part of the source operand. Since the effective addressing mode of the destination register D0, is Data Register Direct, there is no need to retrieve any additional information from memory because all the information that is needed to determine the destination is contained in the opcode word itself. Thus, the complete instruction in memory would be \$303C 0A55. This is illustrated in figure 7.15.

Now, suppose that the instruction uses two absolute addresses for the source effective address and the destination effective address

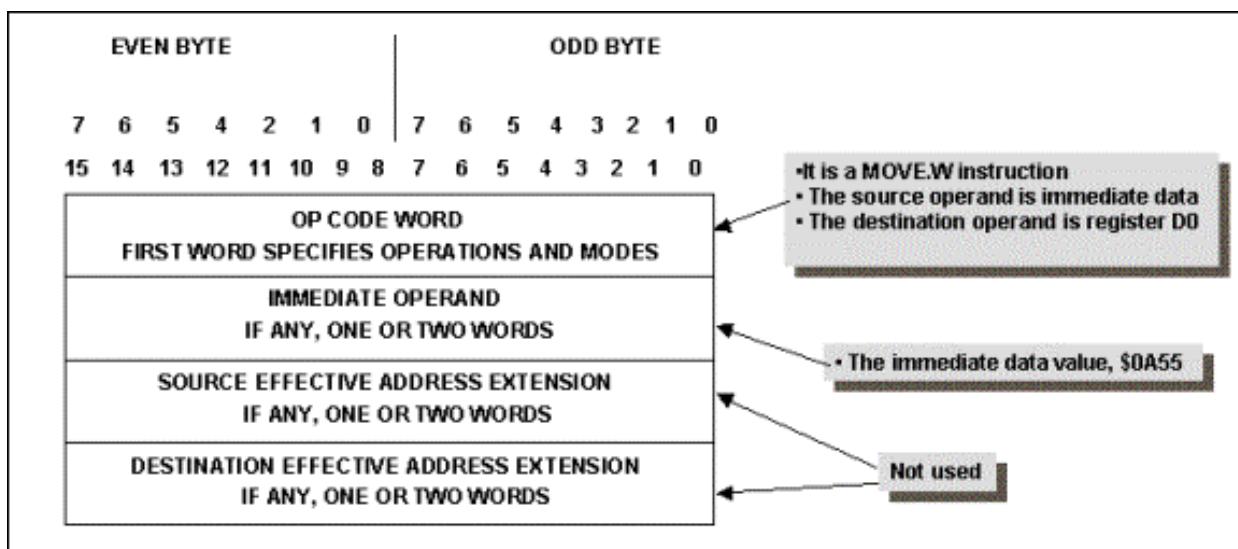


Figure 7.15: Memory storage of the instruction `MOVE.W #$0A55,D0`

`MOVE.W $0A550000,$1000BB00`

This machine language opcode word would decode as 0011 001 111 111 001, or \$33F9. Are we done? Again, the answer is no. This is because there are still two absolute addresses that need to be fetched from memory. The complete instruction looks like

\$33F9 0A55 0000 1000 BB00.

The complete instruction takes up a total of five 16-bit words of memory and requires the processor to make five separate memory-read operations to completely digest the instruction. This is illustrated in figure 7.16.

Word Alignment

The last topic that we need to cover in order to prepare ourselves for the programming tasks ahead of us is that of word alignment in memory. We touched on this topic earlier in the lesson but it is worthwhile to review the concept. Recall that we can address individual bytes in a word by using the least significant address bit, A0, to indicate which

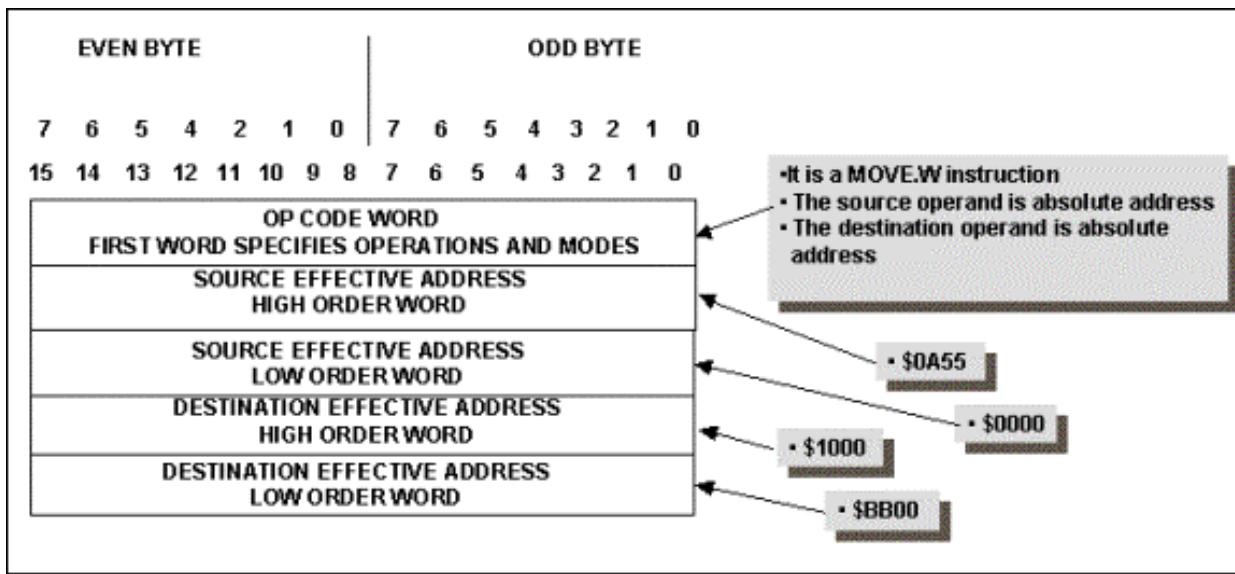


Figure 7.16: Memory storage of the instruction `MOVE.W $0A550000,$1000BB00`

byte we're interested in. However, what would happen if we tried to access a word or a long word value on an odd byte boundary? Figure 7.17 illustrates the problem.

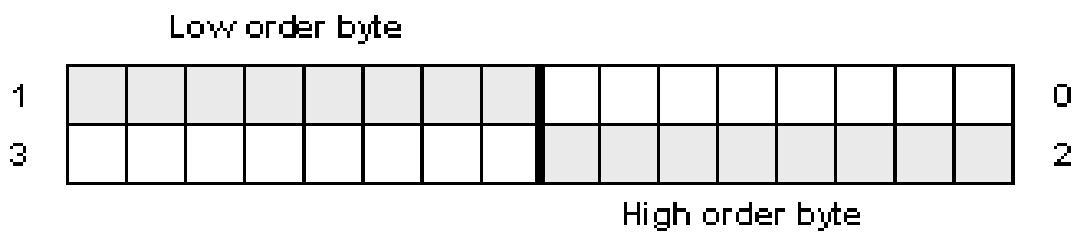


Figure 7.17: Difficulty due to non-aligned access

In order to fetch a word located on an odd-byte boundary the processor would have to make two 16-bit fetches and then discard parts of both words to correctly read the data. Some processors can do this operation. It is an example of the non-aligned access mode that we discussed earlier. In general, this type of access is very costly in terms of processor efficiency. The 68K cannot perform this operation and it will generate an internal exception if it encounters it. The assembler will not catch it. This is a run time error. The cure is to never program a word or long word operation on a byte (odd) memory boundary.

Reading the Programmer's Reference Manual

The most daunting task that you'll face when you set out to write an assembly language program is trying to comprehend the manufacturer's data book. The Motorola 68K *Programmer's Reference Manual* is written in what we could generously call "a rather terse style". It is a reference for people who already know how to write an assembly language program rather than a learning tool for programming students who are trying to learn the methods. From this point on you will need the *Programmer's Reference Manual*, or a good book on 68K assembly language programming, such as *Clements*² textbook. However, the textbook isn't particularly efficient when it comes to writing a program and you'll need a handy reference, such as the "Programmer's Reference Manual." Fortunately, it is easy to get a free copy of the reference manual, either from Freescale, or their corporate website (see the chapter references for the URL).

So, assuming that you have the book in front of you, let's take a quick course in understanding what Motorola's technical writing staff is trying to tell you. The following text is similar to the layout of the reference pages for an instruction.

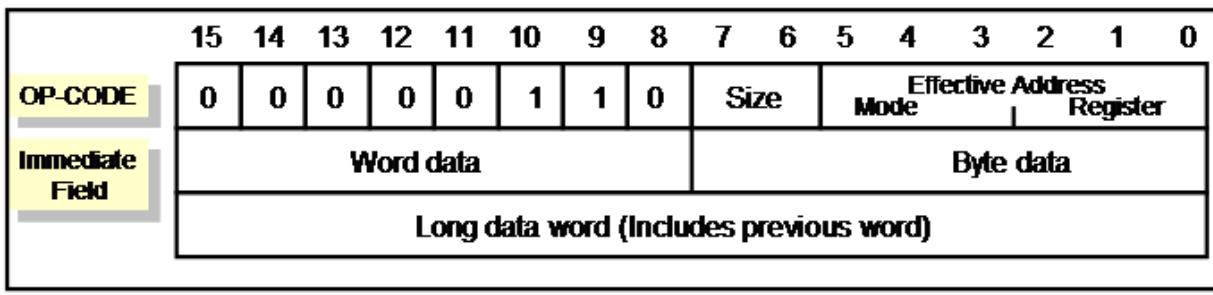
- 1- Heading:** **ADDI** Add Immediate **ADDI**
- 2- Operation:** Immediate Data + Destination → Destination
- 3- Syntax:** ADDI #<data>,<ea>
- 4- Attributes:** Size = (Byte, Word, Long)
- 5- Description:** Add the immediate data to the destination operand and store the result in the destination location. The size of the operation may be specified to be byte, word or long. The size of the immediate data matches the operation size.

6- Condition Codes:

X	N	Z	V	C
●	●	●	●	●

- N Set if the result is negative. Cleared otherwise.
Z Set if the result is zero. Cleared otherwise.
V Set if an overflow is generated. Cleared otherwise.
C Set if a carry is generated. Cleared otherwise.
X Set the same as the carry bit.

7- Instruction format:



Specifies the size of the operation:

DB 7	DB 6	Operation Size
0	0	Byte operation
0	1	Word operation
1	0	Long word operation

9- Effective Address field:

Specifies the destination operand. Only data alterable addressing modes are allowed as shown:

Addr. Mode	Mode	Register	Addr. Mode	Mode	Register
Dn	000	reg. num:Dn	(XXX).W	111	000
An	---	not allowed	(XXX).L	111	001
(An)	010	reg. num:An	#<data>	---	not allowed
(An)+	011	reg. num:An			
-(An)	100	reg. num:An			
(d ₁₆ ,An)	101	reg. num:An	(d ₁₆ ,PC)	---	not allowed
(d ₈ ,An,Xn)	110	reg. num:An	(d ₈ ,PC,Xn)	---	not allowed

10- Immediate field: (Data immediately following the opcode word)

If size = 00, then the data is the low order byte of the immediate word.

If size = 01, then the data is the entire immediate word.

If size = 10, then the data is the next two immediate words..

Let's go through this step-by-step

1.	Instruction and mnemonic: ADDI—Add Immediate
2.	Operation: Immediate Data + Destination → Destination
2.	Add the immediate data to the contents of the destination effective address and place the result back in the destination effective address.

	Assembler syntax: ADDI #(data),<ea>
3.	This explains how the instruction is written in assembly language. Notice that there is a special opcode for an immediate add instruction, ADDI rather than ADD, even though you still must insert the # sign in the source operand field.
4.	Attributes: Byte, word or long word (long) You may add a byte, word or long word. If you omit the .B, .W or .L modifier the default will be .W.
5.	Description: Tells what the instruction does. This is your best hope of trying to understand what the instruction is all about since it is the only real English language on the page.
6.	Condition codes: Lists the condition codes (flags) that are affected by the instruction.
7.	Instruction format: How the machine language instruction is created. Note that there is only one effective address for this instruction.
8.	Effective Address Field: The effective address modes that are permitted for this instruction. It also shows the mode and register values for the effective address. Note that an address register is not allowed as the destination effective address, nor is an immediate data value. Also notice that the two addressing modes involving the Program Counter, PC, are also not allowed to be used with this instruction.

Flow Charting

Assembly language programs, due to their highly structured nature, are good candidates for using flow charts to help plan the structure of a program. In general, flow charts are not used to plan programs written in a high-level language such as C or C++. However, flow charts are still very useful for assembly language program planning. In creating a flow chart, we use a rectangle to represent an "operation." The operation could be one instruction, several instructions, or a subroutine (function call). The rectangle is associated with doing something. We use a diamond to represent a decision point (program flow control), and use arrows to represent program flow. Figure 7.18 is a simple flow chart example of a computer starting up a program.

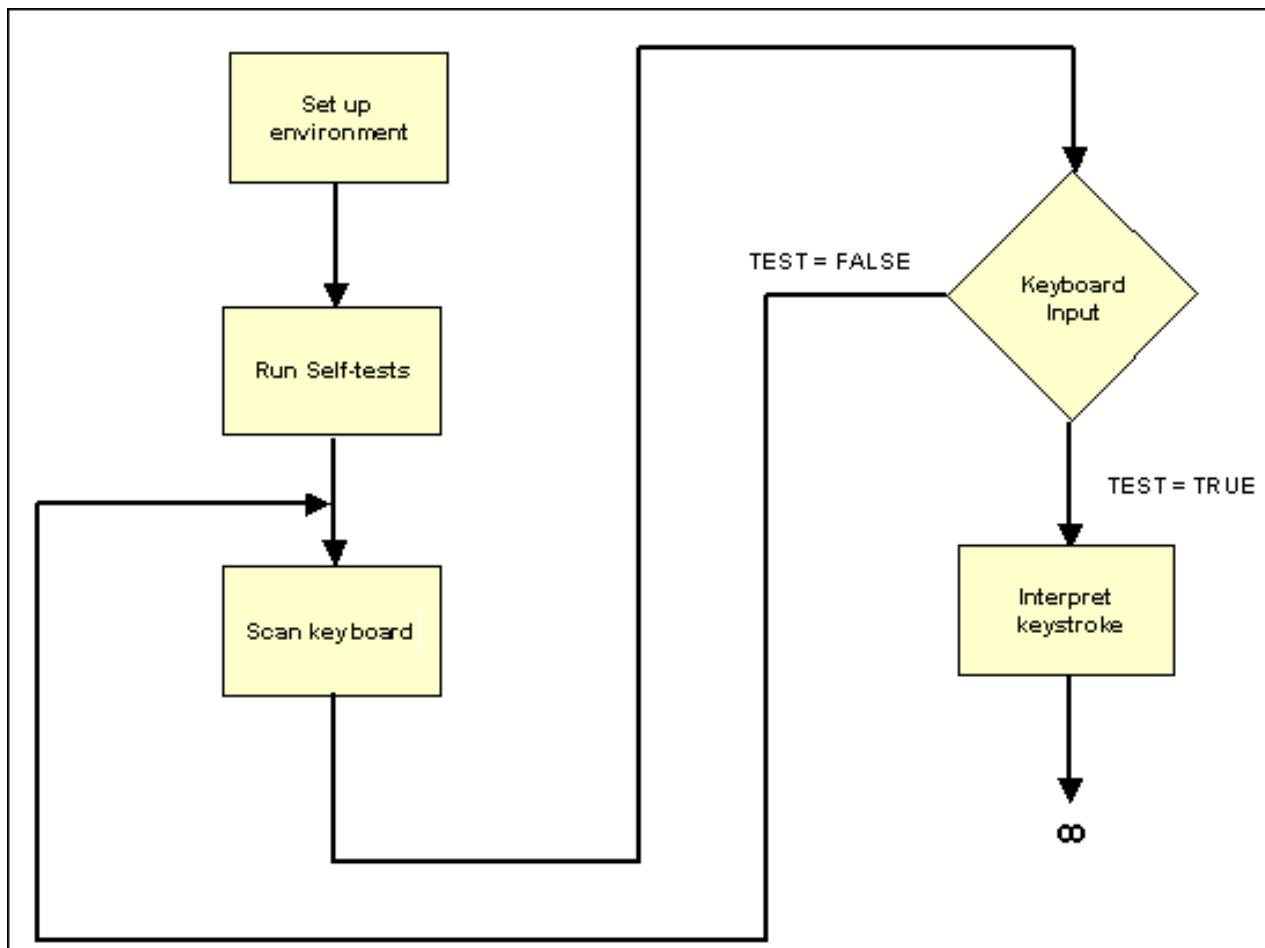


Figure 7.18: Flow chart for a simple program to initialize a system.

The flow chart starts with the operation "Set-up environment." This could be a few instructions or tens of instructions. Once the environment is established, the next operation is to "run the self-tests." Again, this could mean a few instructions or a lot of instructions, we don't know. It depends upon the application. The decision point is waiting for a keyboard input. The program scans the keyboard and then tests to see if a key has been struck. If not, it goes back (loops) and tries again. If a key has been struck, the program then interprets the keystroke and moves on.

Flow-charting is a very powerful tool for helping you to plan your program. Unfortunately, most students (and many professional programmers) take the "code hacking" approach and just jump in and immediately start to write code, much like writing a novel. This probably explains why most programmers wear sandals rather than basketball sneakers with the laces untied, but that's another issue entirely.

Writing an Assembly Language Program

Remember Leonardo DiCaprio's famous line in the motion picture *Titanic*, "I'm king of the world!" In assembly language programming you are the absolute monarch of the computer world. There is nothing that can stop you from making incredible coding blunders. There are no type checks or compiler warnings; just you and the machine, *mano a mano*. In order to write a program in assembly language, you must be continuously aware of the state of the system. You must keep in mind how much memory you have, where it is located, the locations of your peripheral devices, and how to access them. In short, you control everything and you are responsible for everything.

For the purposes of being able to write and execute programs in 68K assembly language, we will not actually use a 68K processor. To do that you would need a computer that was based on a 68K family processor, such as the original Apple Macintosh®. We'll approach it in a different way by using a program called an **instruction set simulator** (ISS) to take the place of the 68K processor.

Commercially available ISSs can cost as much as \$10,000 but ours is free. We're fortunate that two very good simulators have been written for the 68000 processor. The first simulator, developed by Clements⁴ and his colleagues at the University of Teesside, in the United Kingdom, may be downloaded from his website.

Kelley⁵ developed a second simulator, called **Easy68K**. This is a much newer simulator and has extensive debugging support that is lacking in the Teesside version. Both simulators were designed to run under the Windows® operating system. Also, Easy68K was compiled for 32-bit Windows, so it has a much better chance of running under Windows XP and Windows 7^a, although, the Teesside simulator seems to be reasonably well-behaved under the more modern versions of Windows.

The simulators are closer to Integrated Design Environments (IDE). They include a text editor, assembler, simulator and debugger in a package. The simulators are like a computer and debugger combined. You can do many of the debugger operations that you are used to, such as

- peek and poke memory
- examine and modify registers
- set breakpoints
- run from or to breakpoints

^a At the time of this updating of the text, a UWB student, Robert Bartlett-Schneider, is porting Easy68K to the Apple O/S.

- single-step and watch the registers

In general, the steps to create and run an assembly language program are simple and straightforward.

1. Using your favorite ASCII-only text editor, or the one included with the ISS package, create your source program and save it in the old DOS 8.3 file format. Be sure to use the extension .x68. For example, save your program as ***myprog.x68***. The editor that comes with the program will automatically append the x68 suffix to your file.
2. Assemble the program using the included assembler. The assembler will create an absolute machine language file that you may run with the simulator. If your program assembles without errors, you'll then be able to run it in the simulator. The assembler actually creates two files. The absolute binary file and a *listfile*. The listfile shows you your original source program and the hexadecimal machine language file it creates. If there are any errors in your source file, the error will be shown on the listfile output.

In the Teesside assembler, the assembly language program runs on a simulated computer with 1Mbyte of memory occupying the virtual address range of \$00000 . . . \$FFFFF. The ***Easy68K*** simulator runs in a full 16M address space. The Teesside simulator package is not without some minor quirks, but in general, it is well-behaved when running under earlier versions of Windows. However, it may not run under Windows 7, 64-bit, unless you are running a virtual machine environment.

To program in assembly language you should already be a competent programmer. You should have already had several programming classes and understand programming constructs and data structures. Assembly language programming may seem very strange at first, but it is still programming. While C and C++ are free-form languages, assembly language is very structured. Also, it is up to you to keep track of your resources. There is no compiler available to do resource allocation for you.

Don't be overwhelmed by the number of opcodes and operands that are part of the 68K, or any processor's ISA. The real point here is that you can write fairly reasonable programs using just a subset of the possible instructions and addressing modes. Don't be overwhelmed by the number of instructions that you might be able to use. Get comfortable writing programs using a few instructions and addressing modes that you understand and then begin to integrate other instructions and addressing modes when you need something more efficient, or just want to expand your repertoire.

As the programming problems that you will be working on become more involved you'll naturally look for more efficient instructions and effective addressing modes that will make your job easier and result in a better program. As you'll soon discover when we look at different computer architectures, very few programmers or compilers make use of all of the instructions in the instruction set. Most of the time, a small subset of the

instructions will get the job done quite nicely. However, every once in a while, a problem arises that was just made for some obscure instruction to solve. Happy coding.

Pseudo opcodes

There are actually two types of opcodes that you can use in your programs. The first is the set of opcodes that are the actual instructions used by the 68000 processor. They form the 68K ISA. The second set is called *pseudo-ops*, or pseudo opcodes. These are placed in your program just like real opcodes, but they are really instructions to the assembler telling it how to assemble the program. Think of pseudo-ops just as you would think of compiler directives and #define statements.

Pseudo-ops are also called *assembler directives*. They are used to help make the program more readable or to provide additional information to the assembler about how you want the program handled. It is important to realize that commercially available, industrial-strength assemblers are every bit as complex as any modern compiler. Let's look at some of these pseudo opcodes.

- **ORG (set origin):** The **ORG** pseudo-op tells the assembler where in memory to start assembling the program. This is not necessarily where the program might be loaded into memory; it is only telling the assembler where you intend for it to run. If you omit the **ORG** statement, the program will be assembled to run starting at memory location \$00000. Since the memory range from \$00000..\$003FF is reserved for the system vectors, we will generally "ORG" our program to begin at memory address \$00400. Therefore, the first line of your program should be:

```
<label>    ORG      $400    <*comment>
```

The next pseudo-op directive is placed at the end of your source file. It has two functions. First, it tells the assembler to stop assembling at this point and, second, it tells the simulator where to load the program into memory.

- **END (end of source file):** Everything after **END** is ignored. Format:

```
<no label> END    <address>
```

Note that the **ORG** and **END** directives are complementary. The **ORG** directive tells the assembler how to resolve address references; in effect, where you intend for the program to run. The **END** directive instructs the loader program (part of the simulator) where to place the program code in memory. Most of the time the addresses of **ORG** and **END** will

be the same, but they don't have to be. It is quite possible for a program to be loaded in one place and then relocated to another when it is time to run. For our purposes, you would generally start your program with:

```
ORG $400 and end it with: END $400
```

- **EQU (equate directive):** The equate pseudo-op is identical to the `#define` in C. It allows you to provide a symbolic name for a constant value. The format is

```
<label> EQU <expression> <*comment>
```

The expression may be a mathematical expression, but most likely it will just be a number. You may also use the equate directive to create a new symbolic value from other symbolic values. However, the values of the other symbols must be known at the time the new symbol is being evaluated. This means that you cannot have forward references.

The equate directive, like the `"#define"` directive in C and C++, is an instruction to the assembler to substitute the numeric value for the symbolic name in your source file. For example,

```
Bit0_test EQU $01 * Isolate data bit 0  
ANDI.B #Bit0_test,D0 * Is bit 0 in D0 = 1?
```

will be more meaningful to you than:

```
ANDI.B #$01,D0 * Magic numbers
```

especially after you haven't looked at the code for several days.

- **SET (set symbol):** SET is like EQU except that SET may be used to redefine a symbol to another value later on. The format is

```
<label> SET <expression>
```

Data Storage Directives

The next group of pseudo-ops is called *data storage directives*. Their purpose is to instruct the assembler to allocate blocks of memory and perhaps initialize the memory with values.

- **DC (define constant):** Creates a block of memory containing the data values listed in the source file. The format is :

```
<label>           DC.<SIZE>          <item>,<item>,....
```

Example

```
error_msg  DC.B    'Error 99',$0D,$0A,$00  *error message
```

A text string that is placed inside of single quotation marks is interpreted by the assembler as a string of ASCII characters. Thus, this directive is equivalent to writing

```
error_msg  DC.B    $45,$72,$72,$6F,$72,$20,$39,$39,$0D,$0A,$00
```

As you can see, using the single quotes makes your intention much more understandable. In this case, the memory location associated with the label, **error_msg**, will contain the first byte of the string, \$45.

- **DCB (define constant block):** Initialize a block of memory to the same value. The length is the number of bytes, word or long words. The format is

```
<label> DCB.<size>      <length>,<value>
```

- **DS (define storage):** Generates an un-initialized block of memory. Use this if you need to define a storage area that you will later use to store data. The format is:

```
<label> DS.<size>      <length>
```

- **OPT (set options):** Tells the assembler how you want the assembler to create your program code in places where you have not explicitly directed it and on how you want the listfile formatted.

The only option that is worth noting here is the **CRE** option. This option tells the assembler to create a list of cross-references on the listfile. This is an invaluable aid when it comes time to debug your program. For example, examine the block of code in case 1.

CASE 1: Listfile without the CRE option

Source file: EXAMPLE.X68

Defaults: ORG \$0/FORMAT/OPT A,BRL,CEX,CL,FRL,MC,MD,NOMEX,NOPCO

```
1
2 ****
3 *
4         * This is an example of
5         * not using cross-references
6 *
7 ****
8 00000400 ORG $400
9
10 00000400 103C0000      START:      MOVE.B #00,D0
11 00000404 66FA          TEST:       BNE      START
12 00000406 B640          COMPARE:   CMP.W    D0,D3
13 00000408 6BFC          WAIT:      BMI     COMPARE
14 00000400                  END      $400
Lines: 14, Errors: 0, Warnings: 0.
```

CASE 2: Listfile with the CRE option set

Source file: EXAMPLE.X68

Defaults: ORG \$0/FORMAT/OPT A,BRL,CEX,CL,FRL,MC,MD,NOMEX,NOPCO

```
1
2 ****
3 *
4 * This is an example of using cross-
5 * references
6 ****
7
8
9 00000400          OPT  CRE
10 00000400 103C0000  ORG $400
11 00000404 66FA      START:      MOVE.B #00,D0
12 00000406 B640      TEST:       BNE      START
13 00000408 6BFC      COMPARE:   CMP.W    D0,D3
14 00000400          WAIT:      BMI     COMPARE
15 00000400          END      $400
Lines: 15, Errors: 0, Warnings: 0.
```

SYMBOL TABLE INFORMATION

Symbol-name	Type	Value	Decl	Cross reference	line numbers
COMPARE	LABEL	00000406	13	14.	
START	LABEL	00000400	11	12.	
TEST	LABEL	00000404	12	* * NOT USED * *	
WAIT	LABEL	00000408	14	* * NOT USED * *	

Notice how the **OPT CRE** directive creates a symbol table of the labels that you've defined in your program, their values, the line numbers where they are first defined and all the line numbers that refer to them. As you'll soon see, this will become an invaluable aid for your debugging process.

Analysis of an Assembly Language Program

Suppose that we want to implement the simple C assignment statement, $Z = Y + 24$, as an assembly language program. What would the program look like? Let's examine the following simple program. Assume that $Y = 27$. Here's the program.

	ORG	\$400	*Start of code
	MOVE.B	Y,D0	*Get the first operand
	ADDI.B	#24,D0	*Do the addition
	MOVE.B	D0,Z	*Do the assignment
	STOP	#\$2700	*Tell the simulator to stop
	ORG	\$600	*Start of data area
Y	DC.B	27	*Store the constant 27 in memory
Z	DS.B	1	*Reserve a byte for Z
	END	\$400	

Notice that when we use the number 27 without the "\$" preceding it, the assembler interprets it as a decimal number. Also notice the multiple **ORG** statements. This defines a code space at address \$400 and a data space at address \$600. Here's the assembler listfile that we've created. The comments are omitted.

1	00000400	ORG	\$400	
2	00000400	103900000600	MOVE.B	Y,D0
3	00000406	06000018	ADDI.B	#24,D0
4	0000040A	13C000000601	MOVE.B	D0,Z
5	00000410	4E722700	STOP	#\$2700
6	*	Comment line		
7	00000600	ORG	\$600	
8	00000600	1B	Y: DC.B	27
9	00000601	00000001	Z: DS.B	1
10	00000400	END	\$400	

Let's analyze the program line by line:

Line 1: The **ORG** statement defines the starting point for the program.

Line 2: MOVE the byte of data located at address \$600 (Y) into the data register D0. This moves the value 27 (\$1B) into register D0.

Line 3: Adds the byte number 24 (\$18) to the contents of D0 and stores the result back into D0.

Line 4: MOVE the byte of data from D0 to memory location \$601.

Line 5: This instruction is used by the simulator to stop program execution. It is an artifact of the ISS and would not be in most real programs.

Line 6: Comment

Line 7: The **ORG** statement resets the assembler instruction counter to begin counting again at address \$600. This effectively defines the data space for the program. Without the **ORG** statement, the data region would begin immediately after the **STOP** instruction.

Line 8: Defines memory location \$600 with the label "Y" and initializes it to \$1B. Note that even though we use the directive **DC**, there is nothing to stop us from writing to this memory location and changing its value.

Line 9: Defines memory location \$601 with the label "Z" and reserves 1 byte of storage. There is no need to initialize it because it will take on the result of the addition.

Line 10: END directive. The program will load beginning at address \$400.

Notice that our actual program was only three instructions long, but two of the instructions were **MOVE** instructions. This is fairly typical. You'll find that most of your program code will be used to move variables around, rather than do any operations on them.

Summary of Chapter 7

Chapter 7 covered:

- How the memory of a typical computer system may be organized and addressed
- How byte addressing is implemented and the ambiguity of the Big Endian versus Little Endian methods of byte addressing
- The basic organization of assembly language instructions and how the opcode word is interpreted as part of a machine language instruction
- The fundamentals of creating an assembly language program
- The use of pseudo opcodes to control the operation of the assembler program
- Analyzing a simple assembly language program

- Bibliography and references

1- Steven Levy, ***Hackers: Heroes of the Computer Revolution***, ISBN 0-385-19195—2, Anchor Press/Doubleday, Garden City, 1984, pg. 30.

2- Alan Clements, ***68000 Family Assembly Language***, ISBN 0-534-93275-4, PWS Publishing Company, Boston, 1994

3- Motorola Corporation, ***Programmer's Reference Manual***, M68000PM/AD REV 1.
This is also available on-line at:

http://e-www.motorola.com/files/archives/doc/ref_manual/M68000PRM.pdf

4- Alan Clements, <http://www-scm.tees.ac.uk/users/a.clements/>

5- Charles Kelley, <http://www.monroecc.edu/ckelly/tools68000.htm>

Exercises for chapter 7

1- Does the external and internal widths of a processor's data bus have to be the same?
Discuss why they might differ.

2- Explain why the Address bus and Data bus of a microprocessor are called homogeneous and the Status bus is considered to be a heterogeneous bus.

3- All of the following instructions are either illegal or will cause a program to crash. For each instruction, briefly state why the instruction is in error.

a- MOVE.W \$1000,A3

b- ADD.B D0,#\$A369

c- ORI.W #\$55AA007C,D4

d- MOVEA.L D6,A8

e- MOVE.L \$1200F7,D3

4- Briefly explain in a few sentences each of the following terms or concepts.

a- Big Endian/Little Endian:

b- Non-aligned access:

c- Address bus, data bus, status bus:

5- Consider the following 68000 assembly language program. What is the **byte value** stored in memory location \$0000A002 after the program ends?

* System equates

foo	EQU	\$AAAA
bar	EQU	\$5555
mask	EQU	\$FFFF
start	EQU	\$400
memory	EQU	\$0000A000
plus	EQU	\$00000001
magic	EQU	\$2700

* Program starts here

ORG	start	
	MOVE.W	#foo,D0
	MOVE.W	#bar,D7
	MOVEA.L	#memory,A0
	MOVEA.L	A0,A1

```

MOVE.B      D0, (A0)
ADDA.L     #plus,A0
MOVE.B      D7, (A0)
ADDA.L     #plus,A0
MOVE.W      #mask,D3
MOVE.W      D3, (A0)
MOVE.L      (A1),D4
SWAP        D4
MOVE.L      D4,D6
MOVE.L      D6,(A1)
STOP        #magic
END         start

```

6- Examine the code fragment listed below. What is the longword value stored in memory location \$4000?

Start	LEA	\$4000,A0	* Initialize A0
	MOVE.L	#\$AAAAFFFF,D7	* Initialize D7 and D6
	MOVE.L	#\$55550000,D6	
	MOVE.W	D7,D0	* Load D0
	SWAP	D6	* Shell game
	SWAP	D0	
	MOVE.W	D6,D0	* Load D0
	MOVE.L	D0,(A0)	* Save it

7- Examine each of the following Motorola 68,000 assembly language instructions. Indicate which instructions are correct and which are incorrect. For those that are incorrect, write brief explanation of why they are incorrect.

- (a) MOVE.L D0,D7
- (b) MOVE.B D2,#\$4A
- (c) MOVEA.B D3,A4
- (d) MOVE.W A6,D8
- (e) AND.L \$4000,\$55AA

8- Four bytes of data are located in successive memory locations beginning at \$4000. Without using any additional memory locations for temporary storage, reverse the order of the bytes in memory.

9- Examine the following snippet of 68000 assembly language code. What are the **word** contents of D0 after the ADDQ.B instruction? Hint, all logical operations are bit by bit.

```
MOVE.W    #$FFFF, D1
MOVE.W    #$AAAA, D0
EOR.W    D1, D0
ADDQ.B    #01, D0          *What are the word contents of D0?
```

10- Consider the following snippet of assembly code. What is the longword value stored in register D1 at the conclusion of the fourth instruction?

```
START      MOVE.L    #$FA865580, D0
            MOVE.L    D0, $4000
            LSL.W    $4002
            MOVE.L    $4000, D1          * <D1> = ?
```

11- This problem is designed to expose you to the assembler and simulator. At the end of this introduction is a sample program. You should create an assembly source file and then assemble it without any errors. Once it assembles properly, you should then run it in the simulator of your choice. It is best to single-step the simulator from the starting point of the program. The ultimate objective of the exercise is to answer this question, “What is the WORD VALUE of the data in memory location \$4000 when the program is just about to loop back to the beginning and start over again?”

```
*****
*
* My first 68000 Assembly language program
*
*****
* Comment lines begin with an asterisk
* Labels, such as "addr1" and "start", if present, must begin in column
* 1 of a line.
* Opcodes, such as MOVE.W or Pseudo Opcodes, such as EQU, must begin *
in column two or later.
* Watch out for comments, if the text spills over to the next line and
* you forget to use an asterisk, you'll get an assembler error.
*****
*
* Beginning of EQUates section, just like #define in C
*
*****
addr1    EQU    $4000
addr2    EQU    $4001
data2    EQU    $A7FF
data3    EQU    $5555
data4    EQU    $0000
data5    EQU    4678
data6    EQU    %01001111
data7    EQU    %00010111
```

```

*****
*
* Beginning of code segment. This is the actual assembly language
* instructions.
*
*****

```

start	ORG	\$400	* Program starts at \$400
	MOVE.W	#data2,D0	* Load D0
	MOVE.B	#data6,D1	* Load D1
	MOVE.B	#data7,D2	* load D2
	MOVE.W	#data3,D3	* load D3
	MOVEA.W	#addr1,A0	* load address register
	MOVE.B	D1,(A0)+	* transfer byte to memory
	MOVE.B	D2,(A0)+	* transfer second byte
	MOVEA.W	#addr1,A1	* load address
	AND.W	D3,(A1)	* Logical AND

*Stop here. The next instruction shows how a label is used

JMP	start	* Program loops forever
END	\$400	* Stop assembly here

Comments on the program:

- 1- The “EQU”, “END \$400” and “ORG \$400” instructions are the pseudo-op instructions. They are instructions for the assembler, they are not 68000 machine instructions and do not generate any 68000 code.
- 2- Also note that the default number system is decimal. To write a binary number, such as 00110101 you would write it with the percent sign, %00110101. Hexadecimal numbers, such as 72CF, are written with the dollar sign, \$72CF.
- 3- Notice that every instruction is commented.
- 4- Look at the code. Notice that most of the instructions are just moving data around. This is very common in assembly language programming.
- 4- The instruction, **AND.W D3, (A1)** is an example of another address mode called ***indirect*** addressing. In C++ we would call this a pointer. The parentheses around A1 mean that the value contained in the internal register A1 should be considered to be a memory address, and not a data value. This instruction tells the computer to take the 16-bit quantity stored in register D3 and do a logical AND with the 16-bit value stored in memory at the address pointed to by the A1 register and then put the results back into the memory location pointed to by A1.
- 5- The instruction, **MOVE.B D1, (A0) +**, is an example of ***indirect addressing with post increment***. The instruction moves a byte of data (8-bits) from internal data register D1 to the memory location pointed to by the contents of address register A0. In this way it is similar to the previous instruction. However, once the instruction is executed, the contents of A0 are incremented by one byte, so A0 would be pointing to the next byte of memory.